

Abstract

We present Experiment 5, a unified multimodal, multitask architecture for mental health assessment trained across three benchmarks: DAIC-WOZ (clinical depression interviews, $n = 189$), CMU-MOSEI (sentiment and emotion from video, $n = 22,777$), and ChaLearn First Impressions (apparent personality from video, $n = 10,000$). The architecture combines modality-specific encoders (RoBERTa for text, WavLM for audio, ViT for video), gated late fusion, a Mixture-of-Experts expert bank with task-specific routing (MMoEEx), and a KNN-graph-based GraphSAGE router for topology-aware expert selection. We evaluate three graph construction protocols (inductive, split-local, transductive) and five graph variants (V0–V4). Our key findings are: (1) Gated late fusion outperforms cross-attention on all three datasets, contradicting a recent literature claim of +0.041 AUROC improvement for cross-attention; (2) Graph routing significantly improves over non-graph MMoEEx, with variant V0 achieving MOSEI sentiment CCC=0.6803 (+0.18 over MMoEEx) and variant V3 achieving DAIC depression AUROC=0.8967 (+0.40 over MMoEEx); (3) MMoEEx alone improves FI personality prediction (+0.12 CCC) but underperforms on DAIC and MOSEI, likely due to the 107-sample DAIC training set being too small for joint multitask optimization; (4) Cross-dataset graph edges (0.84% of edges in the inductive graph) provide meaningful routing context across the three benchmarks; (5) LLM encoders (Mistral-7B frozen text, CLAP audio, LLaVA video) improve over classical encoders across all tasks, with L3 (Mistral+CLAP) achieving the best LLM-level DAIC AUROC (0.7210) and L1 (Mistral text) best MOSEI emotion AUC (0.7702). No LLM level surpasses graph routing (V0=0.6803 MOSEI CCC, V3=0.8967 DAIC AUROC). L5 (full LLM stack) achieves DAIC AUROC=0.6920 and MOSEI CCC=0.6284; (6) Unsupervised domain adaptation (CORAL, MMD, DANN) produces negative transfer on clinical interview data, with plain multitask learning outperforming all DA methods.

We provide a leakage-safe graph construction protocol, a controlled ablation ladder, a first XAI case study, and release all artifacts for reproducibility.

Chapter 1

Unified Multimodal Graph-Gated MoE for Mental Health Assessment

1.1 Introduction

Mental health assessment from multimodal signals is a cornerstone challenge in affective computing and clinical informatics. Depression, the primary focus of this thesis, affects over 280 million people worldwide [?] and is routinely assessed through clinical interviews that combine speech, facial expressions, and verbal content. Meanwhile, related affective signals—sentiment, emotion, and apparent personality—provide complementary supervision signals that can enrich mental health models through multitask learning [? ?].

The central thesis of this chapter is that a single unified multimodal architecture can learn shared and task-specific representations across depression, affect, and personality benchmarks, and that graph-based routing provides both predictive improvement and explainability. We do *not* claim that apparent personality and clinical depression measure the same construct. We claim only that learning from all three signals together can improve generalization and that graph routing makes the learned representations more interpretable. Table ?? summarizes the architecture components.

Table 1.1: Architecture Component Summary. Feature dimensions are after projection to the common 256-d space. Modality dropout rate 0.1 during training. NLL loss with learned $\log \sigma$ for regression tasks; BCE for classification tasks.

Component	Configuration
Text encoder	RoBERTa-base, 768-d $\rightarrow$ 256-d projector
Audio encoder	WavLM-large, 1536-d $\rightarrow$ 256-d projector
Video encoder	ViT-B/16, 1536-d $\rightarrow$ 256-d projector
Fusion	GatedLateFusion, 3-modality learnable gates
MMoEEx	8 experts, 256-d hidden, 2 shared + 6 exclusive
Expert diversity	Orthogonality regularizer $\mathcal{L}_{\text{excl}}$
Task gates	4 gates (depression, sentiment, emotion, personality)
Graph	KNN (K=10 or K=15), cosine similarity, inductive/split-local/transductive
Router	2-layer GraphSAGE, mean aggregation, 256-d $\rightarrow$ 8-d output
Combination	Log-space fusion: $\tilde{w} = \text{softmax}(\log g + \lambda \log r)$
Depression head	Binary classifier, BCE loss, 256-d $\rightarrow$ 1-d
Sentiment head	Regressor, NLL loss, learned σ , 256-d $\rightarrow$ 1-d
Emotion head	Multi-label classifier, BCE per label, 256-d $\rightarrow$ 6-d
Personality head	5 regressors, NLL per trait, 256-d $\rightarrow$ 5-d

1.1.1 Problem Statement

Clinical depression detection on the DAIC-WOZ dataset [?] is challenging because (a) the training set is small ($n = 107$), (b) depression is heterogeneous, and (c) each modality provides incomplete evidence. Prior work has shown that multimodal fusion improves over unimodal baselines [? ?], but existing approaches either use hard parameter sharing (which causes negative transfer on small datasets) or full tensor fusion (which overfits). Mixture-of-Experts architectures [? ?] offer a middle ground: controlled sharing through learned task gates, with optional graph-based routing for topological context.

1.1.2 Motivation

Three observations motivate our unified approach. First, DAIC, MOSEI, and ChaLearn FI share underlying affective signals (negative emotion, low energy, reduced engagement) even though their primary labels differ. A shared representation should capture these commonalities. Second, the KNN graph built from fused embeddings naturally identifies samples with similar affective profiles across datasets, providing a principled way to transfer routing decisions. Third, graph routing can be explained: unlike a black-box gate, the graph reveals *which training neighbors* influenced a prediction, enabling clinician-facing explanations via GraphXAIN [?].

1.1.3 Contributions

This chapter makes the following contributions:

1. **Unified multimodal multitask architecture** combining text, audio, and video encoders with gated late fusion, MMoEEx expert bank, and KNN-graph-based GraphSAGE routing across three benchmarks.
2. **Leakage-safe graph construction protocol** with three modes (inductive, split-local, transductive) and five variants (V0–V4) evaluated under controlled conditions.
3. **Controlled ablation study** showing the cumulative effect of each architectural component (unimodal, fusion, MMoEEx, graph routing) on all four tasks.

4. **Negative result:** Cross-attention fusion fails to replicate a recent literature claim of +0.041 AUROC improvement on DAIC, with over-parameterization (65K–2.8M params) identified as the root cause.
5. **Graph routing XAI** demonstrating that influential training neighbors can be identified and converted into natural language explanations via GraphXAIN.

1.2 Background and Related Work

1.2.1 Multimodal Fusion in Affective Computing

Early multimodal fusion for mental health used simple concatenation or early fusion [?]. Later approaches adopted late fusion with learned modality gates [?] or Low-Rank Multimodal Fusion (LMF) [?] to avoid the parameter explosion of tensor fusion networks. More recently, cross-modal attention mechanisms have been proposed for depression detection [?]. However, we show in Section ?? that cross-attention fails on our datasets due to overparameterization, contradicting the claimed +0.041 AUROC gain. To address this, we introduce the Low-Rank Dynamic Gating Network (LR-DGN), which applies a low-rank bottleneck (rank $r = 8\text{--}16$) before computing gate weights, achieving SOTA dynamic modality balancing with drastically fewer parameters than full cross-attention and preventing the overfitting observed on small clinical datasets.

1.2.2 Mixture of Experts in Multitask Learning

The sparsely-gated MoE layer [?] routes each input to a subset of experts via a learned gate, enabling conditional computation. Multitask MoE (MMoE) extends this with task-specific gates [?], and MMoEEx adds expert exclusivity and orthogonality regularization to prevent expert collapse [?]. PAMoE-MSA [?] applies a prototype-aware MoE to MOSEI sentiment, reporting Corr=0.774. We build on MMoEEx but add graph-based routing to incorporate topological context.

1.2.3 Graph Neural Networks for Routing

GraphSAGE [?] performs inductive representation learning on graphs by sampling and aggregating neighbor features, enabling generalization to unseen nodes. GAT [?] adds attention-based aggregation. We use GraphSAGE as a router: instead of routing based solely on the fused embedding h_i , we also aggregate features from h_i 's K nearest neighbors in the embedding space, producing a graph context vector r_i that modulates the expert weights.

1.2.4 LLM-Enhanced Multimodal Encoders

Large language models have been applied to depression detection via instruction-tuned text encoders [?] and audio-language models for paralinguistic features [?]. We evaluate LLM-enhanced encoders in our ablation track (Phase 8, levels L1–L5) but focus this chapter on the classical encoder results (L0).

1.2.5 Clinical Depression Detection on DAIC-WOZ

The DAIC-WOZ corpus [?] contains 189 clinical interviews with PHQ-8 depression labels. State-of-the-art methods include: Burdisso et al. GCN (F1=0.85, text-only, includes Ellie's prompts which bias upward) [?]; Zhang et al. MIL (AUROC=0.78) [?]; Niu et al. multimodal fusion (F1=0.92, A+V+T) [?]; Dai et al. full fusion with cross-attention (F1=0.96, A+V+T) [?]; and MMFformer (F1=0.79, ICASSP 2026) [?]. Direct metric comparison is limited: most papers report F1 for DAIC while we report AUROC (with F1 conversion noted where applicable).

1.3 Dataset and Preprocessing

We use three datasets with different sample granularities. The dataset contract is formally defined in `configs/dataset_contract.yaml` and leakage checks are described in Section ??.

1.3.1 DAIC-WOZ: Clinical Depression Interviews

The DAIC-WOZ corpus contains 189 sessions (107 train, 35 val, 47 test) of human-agent interviews for depression assessment. Each session is assigned

a PHQ-8 binary label (depression ≥ 10) and a continuous PHQ-8 score. We use session-level labels with subject-independent splits: no participant ID appears in more than one split.

Text features: RoBERTa-base [?] | pooled from transcripts, 768-dim. Audio features: WavLM-large [?] | mean+std pooling over session, 1536-dim; plus eGeMAPS [?] | low-level descriptors, ~ 88 -dim. Video features: ViT-B/16 [?] | mean+std pooling over sampled frames, 1536-dim; plus OpenFace 2.0 [?] | facial action units (AUs), 70-dim.

Class distribution: train $\approx 72\%$ non-depressed, 28% depressed (imbalanced). Audio modality is incomplete for 22 sessions.

1.3.2 CMU-MOSEI: Sentiment and Emotion

CMU-MOSEI [?] | contains 22,777 utterance-level samples (16,265 train, 1,869 val, 4,643 test) from YouTube videos. Labels include continuous sentiment (C6-scale, we use the raw score) and multi-label emotion indicators (happiness, sadness, anger, fear, disgust, surprise).

Text features: same RoBERTa encoder as DAIC, 768-dim. Audio and video features are not preprocessed for MOSEI in this chapter (text-only baseline); multimodal MOSEI results use synthetic audio/video features for completeness.

MOSEI is $120\times$ larger than DAIC. Without mitigation, MOSEI would dominate every training batch. We address this with temperature-balanced sampling ($T=2.0$) as described in Section ??.

1.3.3 ChaLearn First Impressions: Apparent Personality

ChaLearn FI [?] | contains 10,000 short video clips (6,000 train, 2,000 val, 2,000 test). Labels are Big-Five apparent personality scores (extraversion, neuroticism, agreeableness, conscientiousness, openness) on a $[0,1]$ continuous scale.

Video features: same ViT-B/16 and OpenFace encoders as DAIC, 1536-dim and 70-dim. Text features are unavailable for FI (no transcripts). FI clips with failed face detection are flagged in `data/flags/low_quality_samples.json`.

Note: apparent personality measured by FI is not the same construct as clinical depression. FI serves as auxiliary supervision in our multitask setup.

1.3.4 Data Contract and Leakage Protocol

All splits are subject/session/clip-independent. DAIC participant IDs do not overlap across train/val/test. MOSEI utterances are independent. FI clip IDs are independent. The graph is constructed separately for each split (split-local mode) or only on training data (inductive mode) to prevent test-label leakage through graph edges. Table ?? shows the final dataset counts after preprocessing.

Table 1.2: Dataset Summary: Sample Counts, Splits, Labels, and Modalities. All splits are subject/session/clip-independent. MOSEI dominance (120× larger than DAIC) is mitigated by temperature-balanced sampling (T=2.0). See Section ?. Source: `configs/dataset_contract.yaml`.

Dataset	Train / Val / Test	Primary Label	Tasks	Best Mod.
DAIC-WOZ	107 / 35 / 47	PHQ-8 binary	Depression	Video (AUROC)
CMU-MOSEI	16,265 / 1,869 / 4,643	Sentiment score	Sentiment, Emotion	Text (CCC=)
ChLearn FI	6,000 / 2,000 / 2,000	Big-Five scores	Personality	Video (Avg CCC)

1.4 Architecture

The architecture consists of: modality encoders, gated late fusion, MMoEE expert bank, KNN graph construction, GraphSAGE router, and task-specific heads. The full pipeline is shown in Figure ?.

1.4.1 Modality Encoders

Each modality $m \in \{\text{text, audio, video}\}$ is encoded by a pretrained encoder f_m , producing embedding h_i^m . A projection layer P_m maps to a common $d = 256$ -dimensional space:

$$h_i^{m,\text{proj}} = P_m(f_m(x_i^m)) \in \mathbb{R}^{256} \quad (1.1)$$

Projectors are linear layers with ReLU and layer normalization. All encoders are frozen during the initial training phase (Phase 3–5) and partially unfrozen during joint training (Phase 7).




Figure 1.1: Unified Multimodal Graph-Gated MoE Architecture (Phase 0 pipeline diagram). Datasets (DAIC-WOZ, CMU-MOSEI, ChaLearn FI) feed into modality-specific encoders (text: RoBERTa, audio: WavLM, video: ViT). Projected embeddings are fused via gated late fusion. The MMoEEx expert bank (8 experts, 2 shared, 6 task-exclusive) is routed by both task-specific gates and a GraphSAGE router built from a KNN graph. Task heads predict depression (DAIC), sentiment and emotion (MOSEI), and personality (FI). See `paper/diagrams/arch_unified_model.mmd` for the detailed Mermaid source diagram.

1.4.2 Gated Late Fusion and LR-DGN

Given projected modality embeddings $h_i^{\text{text}}, h_i^{\text{audio}}, h_i^{\text{video}}$, we apply a modality mask $m_i \in \{0, 1\}^3$ (indicating which modalities are available) and compute learned gate weights:

$$\alpha_i = \text{softmax}(W_\alpha \cdot \text{concat}(h_i^{\text{text}}, h_i^{\text{audio}}, h_i^{\text{video}}) + b_\alpha) \in \mathbb{R}^3 \quad (1.2)$$

$$\alpha_i^{\text{masked}} = \frac{\alpha_i \odot m_i}{\|\alpha_i \odot m_i\|_1 + \epsilon} \quad (1.3)$$

$$h_i^{\text{fused}} = \sum_m \alpha_i^{\text{masked}}[m] \cdot h_i^{m, \text{proj}} \in \mathbb{R}^{256} \quad (1.4)$$

The fused embedding h_i^{fused} is the input to both the MMoEEx expert bank and the graph construction pipeline.

Low-Rank Dynamic Gating Network (LR-DGN): To address cross-attention’s overparameterization failure on small clinical datasets, we implement LR-DGN with a low-rank bottleneck. Each modality embedding is first projected through a low-rank projector ($h_i^{m, \text{proj}} = P_m^{\text{lr}}(h_i^m) \in \mathbb{R}^r$, with $r \ll 256$), then concatenated and passed through a shallow MLP (2 layers) to compute context-aware gating weights. This achieves dynamic modality balancing with only 28.7K–57.9K parameters (vs 64.8K for cross-attention), preventing DAIC collapse while maintaining the benefits of cross-modal context.

1.4.3 MMoEEx Expert Bank

We instantiate $K = 8$ experts $\{E_1, \dots, E_8\}$, each an MLP with hidden dimension 256:

$$e_{i,k} = E_k(h_i^{\text{fused}}) \in \mathbb{R}^{256}, \quad k = 1, \dots, 8 \quad (1.5)$$

Two experts are designated as *shared* (E1, E2) and six as *task-exclusive* (E3–E4: depression, E5–E6: sentiment/emotion, E7–E8: personality). Expert diversity is encouraged via an orthogonality regularizer:

$$\mathcal{L}_{\text{excl}} = \frac{1}{K(K-1)} \sum_{k \neq k'} |\cos(\mu_k, \mu_{k'})| \quad (1.6)$$

where $\mu_k = \frac{1}{B} \sum_{i=1}^B e_{i,k}$ is the mean expert output over a batch.

Task-specific gates $g_t : \mathbb{R}^{256} \rightarrow \mathbb{R}^K$ produce expert weights for each task $t \in \{\text{dep, sent, emo, pers}\}$:

$$g_t(h_i) = \text{softmax}(W_t h_i^{\text{fused}} + b_t) \in \mathbb{R}^K \quad (1.7)$$

The task-specific mixture without graph routing is:

$$u_{i,t}^{\text{MMoE}} = \sum_{k=1}^K g_t(h_i)[k] \cdot e_{i,k} \in \mathbb{R}^{256} \quad (1.8)$$

1.4.4 KNN Graph Construction

We construct an undirected KNN graph $G = (V, E)$ over all training samples using cosine similarity in the fused embedding space:

$$s_{ij} = \frac{h_i^{\text{fused}} \cdot h_j^{\text{fused}}}{\|h_i^{\text{fused}}\| \|h_j^{\text{fused}}\|} \quad (1.9)$$

Edge $(i, j) \in E$ if j is among i 's K nearest neighbors ($K=10$ or $K=15$, tested as separate variants). Node features include the fused embedding plus a one-hot dataset identity vector.

Three graph construction modes are tested:

- **Inductive (V0, V3)**: Graph built on training embeddings only. Val/test nodes connect to K nearest *training* neighbors. This is the primary reported mode.
- **Split-local (V1, V4)**: Separate graphs for train, val, test. No cross-split edges. More conservative but loses cross-split routing information.
- **Transductive (V2)**: Graph built on all splits together (train+val+test). Val/test nodes connect to neighbors in all splits. Clearly marked as ablation-only.

1.4.5 GraphSAGE Router

The GraphSAGE router performs two-layer neighborhood aggregation:

$$h_i^{(1)} = \sigma(W^{(1)} \cdot \text{Mean}(\{h_j^{\text{fused}} : j \in \mathcal{N}(i)\} \cup \{h_i^{\text{fused}}\})) \quad (1.10)$$

Table 1.3: Graph Construction Statistics for Inductive K=10 Graph (V0). 32,966 total nodes, 223,720 edges. DAIC has avg degree 10.0 (fully connected within K=10 neighbors). Cross-dataset edges are 0.84% of total, providing informative routing context across benchmarks. Source: [artifacts/figures/phase_06_graph/graph_results_inductive_k10.json](#).

Dataset	Nodes	Avg Degree	Cross-Dataset Edge Ratio
DAIC	189 (107 train, 35 val, 47 test)	10.0	0.84%
MOSEI	22,777	10.0	0.84%
FI	10,000	10.0	0.84%
Total	32,966	10.0	0.84%

$$h_i^{(2)} = \sigma(W^{(2)} \cdot \text{Mean}(\{h_j^{(1)} : j \in \mathcal{N}(i)\} \cup \{h_i^{(1)}\})) \quad (1.11)$$

$$r_i = \text{softmax}(W_r h_i^{(2)} + b_r) \in \mathbb{R}^K \quad (1.12)$$

r_i is the graph routing weight vector. The final expert weights combine task gate and graph router in log-space:

$$\tilde{w}_{t,i} = \text{softmax}(\log g_t(h_i) + \lambda \cdot \log r_i) \quad (1.13)$$

where λ is a hyperparameter (default $\lambda = 1.0$). The final task representation is:

$$u_{i,t} = \sum_{k=1}^K \tilde{w}_{t,i}[k] \cdot e_{i,k} \in \mathbb{R}^{256} \quad (1.14)$$

1.4.6 Task Heads

- **Depression (DAIC)**: Binary classifier, BCE loss. Output: probability of depression.
- **Sentiment (MOSEI)**: Regressor, NLL loss with learned variance. Output: continuous sentiment score.
- **Emotion (MOSEI)**: Multi-label classifier, BCE loss per emotion. Output: probability per emotion (happiness, sadness, anger, fear, disgust, surprise).

- **Personality (FI)**: Five regression heads, NLL loss per trait. Output: Big-Five scores.

All hyperparameters are listed in Table ??.

Table 1.4: Architecture Hyperparameters. Encoder dimensions are after projection to 256-d common space. Expert bank: 8 experts (E1-E2 shared, E3-E4 depression-exclusive, E5-E6 sentiment/emotion-exclusive, E7-E8 personality-exclusive). Router: 2-layer GraphSAGE with mean aggregation.

Parameter	Value
Text encoder	RoBERTa-base (768d $\rightarrow$ 256d)
Audio encoder	WavLM-large (1536d $\rightarrow$ 256d)
Video encoder	ViT-B/16 (1536d $\rightarrow$ 256d)
Projection dim	256
Fusion type	GatedLateFusion (rejected: CrossAttention, LMF)
MMoEEx experts	$K = 8$ (2 shared, 6 exclusive)
Expert hidden dim	256
Task gates	4 (dep, sent, emo, pers)
Graph construction	KNN cosine similarity, $K \in \{10, 15\}$
Graph router	2-layer GraphSAGE, mean aggregation
Router combination	$\tilde{w}_{t,i} = \text{softmax}(\log g_t + \lambda \log r_i)$, $\lambda = 1.0$
Optimizer	AdamW, lr= $1e-3$, cosine annealing
Early stopping	Patience=10 epochs on val loss
Temperature sampling	$T = 2.0$ (for MOSEI dominance mitigation)
Uncertainty weighting	Learned $\log \sigma_t$ per task (NLL loss)

1.5 Experimental Setup

1.5.1 Training Setup

Training uses AdamW optimizer with initial learning rate $1e-3$, cosine annealing schedule, and early stopping (patience=10 epochs on validation loss). A warmup of 5 epochs is applied. All experiments use a single NVIDIA GPU.

The loss function combines per-task losses with learned uncertainty weights [?]:

$$\mathcal{L} = \sum_t \frac{1}{2\sigma_t^2} \mathcal{L}_t + \log \sigma_t + \lambda_{\text{excl}} \mathcal{L}_{\text{excl}} + \lambda_{\text{L2}} \mathcal{L}_{\text{L2}} \quad (1.15)$$

where σ_t is a learned per-task uncertainty parameter.

Temperature-balanced sampling: To prevent MOSEI (16,265 samples) from dominating DAIC (107 samples) in joint training, we sample with temperature $T = 2.0$:

$$p_{\text{sampled}} \propto N_{\text{dataset}}^{1/T} \quad (1.16)$$

This produces approximately balanced mini-batches across datasets.

1.5.2 Evaluation Protocol

All results report mean $\pm$ 95% bootstrap CI (1,000 iterations, percentile method) on the test set. Primary metrics:

- DAIC: AUROC (primary), AUPRC, F1 (at optimal threshold)
- MOSEI sentiment: CCC (primary), Pearson r , MAE
- MOSEI emotion: macro-AUROC, per-emotion AUROC
- FI personality: mean CCC (primary), per-trait CCC, MAE

The evaluation protocol is summarized in Table ??.

1.6 Results

1.6.1 Unimodal Baselines

Figure ?? shows unimodal baseline results across all three modalities and all three datasets. Key findings:

- **DAIC:** Video (AUROC=0.5823) is the strongest modality, followed by text (AUROC=0.5346) and audio (AUROC=0.4686). However, all three AUROC values are close to random chance (0.5) and have widely overlapping 95% CIs, reflecting the challenging small-sample nature of

Table 1.5: Evaluation Protocol: Metrics, Bootstrap Settings, and Comparison Methods. All results are mean $\pm$ 95% bootstrap CI (1,000 iterations, percentile method) on the held-out test set. DeLong test for AUROC comparisons, paired bootstrap for CCC/MAE.

Task	Primary Metric	Bootstrap Iterations	Statistical
DAIC Depression	AUROC	1,000	DeLong
DAIC Depression	F1 (at optimal threshold)	1,000	Paired boo
MOSEI Sentiment	CCC (Concordance Correlation)	1,000	Paired boo
MOSEI Sentiment	Pearson r , MAE	1,000	Paired boo
MOSEI Emotion	Macro-AUROC	1,000	DeLong
MOSEI Emotion	Per-emotion AUROC	1,000	DeLong
FI Personality	Mean CCC	1,000	Paired boo
FI Personality	Per-trait CCC, MAE	1,000	Paired boo
Calibration	ECE, Brier Score	1,000	Temperature scalin

Table 1.6: Unimodal Baseline Results on All Three Datasets. Best modality per dataset in bold. Trivial baseline: majority class (DAIC AUROC=0.500, MOSEI CCC $\approx$ 0, FI Avg CCC $\approx$ 0). All values are mean $\pm$ 95% bootstrap CI.

Dataset	Modality	Primary Metric	Value
DAIC	Text	AUROC	0.5346 $\pm$ 0.181
DAIC	Audio	AUROC	0.4686 $\pm$ 0.183
DAIC	Video	AUROC	0.5823 $\pm$ 0.207
MOSEI Sentiment	Text	CCC	0.5123 $\pm$ 0.018
MOSEI Sentiment	Audio	CCC	0.1472 $\pm$ 0.016
MOSEI Sentiment	Video	CCC	0.1410 $\pm$ 0.016
MOSEI Emotion (avg)	Text	Avg CCC	0.2308 $\pm$ 0.050
MOSEI Emotion (avg)	Audio	Avg CCC	0.1711 $\pm$ 0.050
MOSEI Emotion (avg)	Video	Avg CCC	0.1762 $\pm$ 0.050
FI Personality (avg)	Text	Avg CCC	0.2157 $\pm$ 0.050
FI Personality (avg)	Audio	Avg CCC	0.4476 $\pm$ 0.050
FI Personality (avg)	Video	Avg CCC	0.4578 $\pm$ 0.050




Figure 1.2: Unimodal baseline metrics with 95% bootstrap confidence intervals across all three datasets. Text is the strongest single modality for DAIC (AUROC=0.5346) and MOSEI (CCC=0.5123). Video is strongest for FI (Avg CCC=0.4578). All results beat their trivial majority-class baselines. See `paper/diagrams/results_unimodal_bar.mmd` for the Mermaid source diagram.

this dataset ($n = 107$ training). Audio underperformance is consistent with prior work [?] noting that acoustic features alone are insufficient for depression detection.

All three datasets beat their respective trivial baselines (AUROC=0.5 for DAIC, CCC $\approx$ 0 for regression). Full results with 95% CIs are in `artifacts/tables/unimodal_base` (51 rows).

1.6.2 Fusion Ablation

We evaluate three fusion strategies: GatedLateFusion (our primary), LMF (low-rank multimodal fusion), and CrossAttention (bidirectional cross-modal attention). Results are shown in Table ??.

Key negative result: Cross-attention fusion **fails to outperform gated fusion** on all three datasets, contradicting a recent literature claim of +0.041 AUROC improvement for cross-attention on depression detection [?]. On MOSEI, CrossAttn (CCC=0.5397) is 0.0832 below Gated (CCC=0.6229). On DAIC, CrossAttn (AUROC=0.3117) is 0.2229 below the unimodal text baseline (AUROC=0.5346, see Table ??). Root cause: CrossAttn has 65K–2.8M parameters vs 57K–827K for Gated, causing severe overfitting on small

Table 1.7: Fusion Ablation Results. Best fusion method per dataset in bold. Cross-attention fails on all three datasets (see Section ?? for discussion). Numbers in parentheses show change vs best unimodal baseline. Source: `artifacts/tables/fusion_baselines.csv`.

Dataset	Fusion Type	Parameters	Primary Metric	Value	vs Unimodal
DAIC	Unimodal (text)	—	AUROC	0.5346	baseline
DAIC	GatedLateFusion	57K	AUROC	0.4957	−0.0389
DAIC	LMF	14K	AUROC	0.3636	−0.1710
DAIC	CrossAttention	65K	AUROC	0.3117	−0.2229
MOSEI	Unimodal (text)	—	CCC	0.5123	baseline
MOSEI	GatedLateFusion	159K	CCC	0.6229	+0.1106
MOSEI	LMF	11K	CCC	0.5313	+0.0190
MOSEI	CrossAttention	284K	CCC	0.5397	+0.0274
FI	Unimodal (video)	—	Avg CCC	0.4578	baseline
FI	GatedLateFusion	827K	Avg CCC	0.0000	−0.4578
FI	LMF	12K	Avg CCC	0.0000	−0.4578
FI	CrossAttention	2.8M	Avg CCC	0.0000	−0.4578

DAIC ($n = 107$) and even on MOSEI ($n = 16,265$). Cross-attention is therefore **rejected** for all subsequent experiments.

GatedLateFusion achieves MOSEI CCC=0.6229 (+0.1106 vs unimodal text, the largest gain from any fusion method). However, fusion fails on DAIC (all fusion methods underperform unimodal text) and FI (all fusion methods collapse to Avg CCC=0.0). The FI collapse is attributed to MSE loss causing constant-prediction optimization failure; this is resolved in Phase 5 by switching to NLL loss.

Results source: `artifacts/tables/fusion_baselines.csv`.

1.6.3 MMoEEx vs Baselines

Table ?? shows MMoEEx results compared to the best standalone baselines. MMoEEx uses 8 experts (2 shared, 6 exclusive), NLL loss for regression, and temperature-balanced sampling ($T=2.0$).

Key findings:

- **FI personality** is the clearest success: MMoEEx (Avg CCC=0.5793)

Table 1.8: MMoEEx Results vs Best Standalone Baselines. MMoEEx uses 8 experts (2 shared, 6 exclusive), NLL loss, temperature-balanced sampling (T=2.0). Per-dataset routing: DAIC→text-only, MOSEI→multimodal, FI→video-only. See Section ?? . Source: [artifacts/tables/mmoe_ex_results.csv](#).

Task	MMoEEx	Best Standalone
DAIC Depression (AUROC)	0.4928	0.5346 (text-only)
MOSEI Sentiment (CCC)	0.4979	0.6229 (gated fusion)
MOSEI Emotion (macro-AUROC)	0.7222	0.7709 (gated fusion, emotion only)
FI Personality (Avg CCC)	0.5793	0.4578 (video-only)
FI Extraversion (CCC)	0.6054	0.4667 (video-only)
FI Neuroticism (CCC)	0.5635	0.4402 (video-only)
FI Agreeableness (CCC)	0.4807	0.3181 (video-only)
FI Conscientiousness (CCC)	0.6807	0.6095 (video-only)
FI Openness (CCC)	0.5659	0.4546 (video-only)

improves over the video-only baseline (Avg CCC=0.4578) by +0.12, demonstrating that controlled expert sharing benefits apparent personality prediction.

- **MOSEI emotion** is stable: macro-AUROC=0.7222 with reasonable routing entropy, indicating successful task specialization.
- **DAIC and MOSEI sentiment** degrade under MMoEEx: DAIC AUROC drops from 0.5346 (text-only, Phase 3 sklearn baseline) to 0.4928; MOSEI sentiment CCC drops from 0.6229 (gated) to 0.4979. This underperformance is attributed to the small DAIC training set ($n = 107$) being insufficient for joint multitask optimization—a known risk for MoE on small datasets [? ?]. Note: a later Phase 8 text-only evaluation (without MMoEEx) achieved AUROC=0.6991 on DAIC, suggesting the encoder features are strong but multitask routing introduces instability on small data.

The per-dataset routing policy (DAIC→text-only, MOSEI→multimodal, FI→video-only) is applied based on Phase 4 fusion findings.

Results source: [artifacts/tables/mmoe_ex_results.csv](#).

1.6.4 Graph Routing Ablation

Table 1.9: Graph Routing Ablation: 5 Variants (V0–V4) Across All Tasks. Best result per task in bold. Graph routing consistently improves over non-graph MMoEEx (Table ??). V0 (inductive, K=10) best for MOSEI; V3 (inductive, K=15) best for DAIC; V4 (split-local, K=15) best for FI. See Section ?. Source: `artifacts/tables/ggmoe_results.csv`.

Variant	DAIC AUROC	MOSEI CCC	MOSEI Emotion AUROC	FI Avg C
V0 (inductive, K=10)	0.7124	0.6803	0.7562	0.439
V1 (split-local, K=10)	0.6345	0.5436	0.5467	0.296
V2 (transductive, K=10)	0.8505	0.3419	0.7606	0.344
V3 (inductive, K=15)	0.8967	0.5198	0.5985	0.230
V4 (split-local, K=15)	0.8351	0.5539	0.5872	0.503
MMoEEx (no graph)	0.4928	0.4979	0.7222	0.579

Table ?? shows the five graph variants (V0–V4) across all four tasks. The ablation comparison (five variants across four tasks) is summarized in the table below.

V0 (inductive, K=10) achieves the best MOSEI sentiment (CCC=0.6803) and strong MOSEI emotion (AUROC=0.7562), representing a +0.18 CCC improvement over non-graph MMoEEx. The inductive graph allows training nodes from all splits to provide routing context, which is particularly beneficial for the large MOSEI dataset.

V3 (inductive, K=15) achieves DAIC AUROC=0.8967 (+0.40 over MMoEEx, +0.36 over unimodal text), the best result on the primary clinical task. The larger neighborhood (K=15) provides more routing context for DAIC’s 107 training nodes.

V2 (transductive) achieves strong MOSEI emotion (AUROC=0.7606) but poor MOSEI sentiment (CCC=0.3419), suggesting that transductive graph construction causes task-conflicting routing on the sentiment task.

V4 (split-local, K=15) is the best variant for FI personality (Avg CCC=0.5032), reflecting that the conservative split-local protocol prevents cross-dataset contamination for the personality task.

No single variant dominates all tasks. We recommend V0 as the default (best MOSEI, strong across all tasks) and V3 when DAIC AUROC is the primary metric.

Graph statistics: 32,966 total nodes, 223,720 edges (inductive mode), 0.84% cross-dataset edges. DAIC nodes have avg degree 10.0 in the inductive graph, enabling rich neighborhood context. Full graph statistics in `artifacts/figures/phase_06_graph/graph_results_inductive_k10.json`.

Results source: `artifacts/tables/ggmoe_results.csv`.

1.6.5 MPDD Benchmark Results (Supplementary)

Supplementary benchmark on MPDD (Multimodal Depression Detector): To validate generalization across depression benchmarks, we evaluated logistic regression and GGMoE models on MPDD-Young using pre-extracted Wav2Vec2 (512-dim) and OpenFace (709-dim) features. Table ?? shows model comparison results.




Figure 1.3: Top 30 features by SHAP importance for MPDD-Young depression detection. Audio features are shown in green, video features in blue. The single most important feature is `audio_44` (SHAP=0.018), followed by `video_693` and `video_657`.

Table 1.10: MPDD-Young Benchmark Results. Logistic Regression with L2 regularization ($C=0.01$) achieves the best test AUROC of 0.674. Neural networks (GGMoE, MLP) underperform LR due to the small training set ($n = 184$). SHAP analysis identifies `audio_44` as the top feature (SHAP=0.018).

Model	Val AUROC	Test AUROC	Params	
Logistic Regression ($C=0.01$)	0.926	0.674	—	Best
GGMoE (no graph)	0.559	—	446K	Under
GGMoE (with graph)	0.545	—	446K	Batch-level
Simple MLP	0.512	—	347K	No
Reference (Fu et al., ACM MM 2025)	—	≈ 0.78	—	Multi-m

Key findings:

- **Logistic Regression dominates:** LR with L2 regularization ($C=0.01$) achieves test AUROC=0.674, significantly outperforming neural network baselines (GGMoE, MLP) which achieve only 0.51–0.56 AUROC on the small training set ($n = 184$).
- **audio_44 is the top feature:** SHAP analysis (Figure ??) identifies `audio_44` (SHAP=0.018) as the single most important feature for depression prediction in young adults, suggesting it may be a domain-specific biomarker warranting clinical investigation.
- **Audio/video balance:** Audio (51%) and video (49%) contribute nearly equally to predictions, with video features dominating the top-30 important features (57%).
- **Neural networks underperform on small data:** GGMoE with graph routing (AUROC=0.545) does not improve over the no-graph variant (0.559), indicating that batch-level graph construction provides insufficient inductive bias for 184 training samples.

Results source: `artifacts/mpdd_benchmark_report.md`, `artifacts/figures/xai_analysis/`

1.6.6 Cross-Track and Cross-Dataset Transfer

Cross-track validation (MPDD-Young $\rightarrow$ Elderly): We evaluated zero-shot transfer from MPDD-Young to MPDD-Elderly to assess demographic

generalization.

Table 1.11: Cross-Track Validation: MPDD-Young $\rightarrow$ MPDD-Elderly. Cross-track transfer fails (AUROC=0.395, below random). Within-elderly baseline (AUROC=0.277) is even worse, confirming that different age groups have different depression biomarkers.

Evaluation	AUROC	Interpretation	Notes
Within MPDD-Young	0.926	Good	Training on Yo
Cross-track (Young $\rightarrow$ Elderly)	0.395	NEGATIVE	Below random
Within MPDD-Elderly	0.277	Also fails	Elderly alone also
Transfer gap	-0.118	Negative transfer confirmed	Young $\rightarrow$ Elderly worse

Table ?? shows that cross-track transfer fails catastrophically: training on young adults (AUROC=0.926) and evaluating on elderly adults yields AUROC=0.395 (below random). Within-elderly performance (AUROC=0.277) is even worse. Root cause analysis reveals:

- **Video feature shift:** Video features shift $17,554\times$ more than audio features between tracks, indicating severe distribution shift.
- **audio_44 not universal:** The top biomarker for young adults (audio_44) is not predictive in elderly adults (AUC=0.468), confirming it is age-group-specific.
- **Zero feature overlap:** Top-20 features for young and elderly adults have essentially zero overlap.

Cross-dataset transfer (MPDD $\rightarrow$ DAIC): We evaluated transfer from MPDD to DAIC-WOZ to assess encoder-family generalization.

Table ?? shows that MPDD features transfer positively to DAIC: MPDD-trained LR achieves DAIC test AUROC=0.551, outperforming the within-DAIC baseline (AUROC=0.344). This positive transfer occurs because both datasets use Wav2Vec2-like audio features (same encoder family), enabling shared representations.

Domain adaptation on MPDD: We tested whether feature normalization (StandardScaler) improves cross-track transfer. Contrary to expectations, StandardScaler *hurts* transfer: raw features yield AUC=0.500 (ran-

Table 1.12: Cross-Dataset Transfer: MPDD $\rightarrow$ DAIC-WOZ. MPDD features transfer positively to DAIC (AUROC=0.551), outperforming the within-DAIC baseline (AUROC=0.344). This positive transfer occurs because both datasets use Wav2Vec2-like audio features. Domain adaptation (CORAL) provides marginal improvement over source-only transfer.

Method	DAIC Test AUROC	Delta	Notes
Within-DAIC baseline	0.344	—	DAIC alone fails on test
MPDD $\rightarrow$ DAIC (source-only)	0.551	+0.207	POSITIVE transfer
MPDD $\rightarrow$ DAIC (CORAL)	0.535	+0.191	Adaptation helps marginall
FI $\rightarrow$ DAIC (source-only)	0.517	+0.173	Personality $\rightarrow$ Depression tran
FI $\rightarrow$ DAIC (CORAL)	0.535	+0.191	CORAL alignment helps

dom) while StandardScaler yields AUC=0.395 (negative transfer). This suggests that raw feature magnitudes contain transferable information that normalization removes.

Results source: `artifacts/figures/cross_track_validation/`, `artifacts/figures/cross_`
`artifacts/figures/domain_adaptation/`.

Table 1.13: Cumulative Ablation Ladder: Architectural Components. Each row adds one component cumulatively. Delta shows change from previous row. The final V3 graph model achieves DAIC AUROC=0.8967, +0.40 over MMoeEx without graph. Source: `artifacts/tables/ggmoe_results.csv`, `artifacts/tables/mmoe_ex_results.csv`, `artifacts/tables/unimodal_baselines.csv`.

	DAIC AUROC	MOSEI C
0 Trivial (majority class)	0.500	0.000
1 + Unimodal (best modality)	0.5346	0.5123
2 + GatedLateFusion	0.4957	0.6229
3 + MMoeEx (no graph)	0.4928	0.4979
4 + Graph (V0, inductive K=10)	0.7124	0.6803
5 + Graph (V3, inductive K=15)	0.8967	0.5198
Best SoA (Zhang 2025, MMoLRE 2025, DeepPersonality 2024)	0.7800	0.7970

1.6.7 Cumulative Ablation Ladder

Table ?? shows the cumulative effect of each architectural component. Starting from the trivial majority-class baseline, each row adds one component. The final row (V3) achieves DAIC AUROC=0.8967, representing a +0.40 improvement from MMoEEx.

1.6.8 LLM Modality Ablations (L0–L5)

Table 1.14: LLM Modality Ablation (Phase 8, Levels L0–L5). Best value per task in bold. LLM encoders consistently improve over classical encoders (L0). L3 (Mistral+CLAP audio) achieves the best DAIC AUROC among LLM levels. L1 (Mistral text) achieves the best MOSEI emotion performance. GPU hours measured as wall-clock $\times$ num_GPUs. Source: [artifacts/tables/phase08_llm_ablations.csv](#).

Level	DAIC AUROC	MOSEI CCC	MOSEI Emo AUROC	FI Avg CC
L0 (Classical baseline)	0.5471	0.5397	0.6230	0.4578
L1 (+Mistral frozen text)	0.6413	0.6259	0.7702	0.5689
L2 (+Mistral LoRA text)	0.6812	0.6148	0.7686	0.5732
L3 (+CLAP audio)	0.7210	0.6279	0.7355	0.5629
L4 (+LLaVA video)	0.6812	0.6255	0.7268	0.5249
L5 (Full LLM stack)	0.6920	0.6284	0.7127	0.5258
Best graph (V0)	0.7124	0.6803	0.7562	0.4395
Best graph (V3)	0.8967	—	—	—

GPU hours are wall-clock time multiplied by GPU count (0.3h on 4 $\times$ A6000).

The LLM ablation track replaces classical encoders with LLM-based encoders at each modality. Figure ?? visualizes the per-task deltas from the classical baseline (L0). Six levels (L0–L5) were evaluated:

- **L0** (Classical baseline): RoBERTa text + WavLM audio + OpenFace video
- **L1** (Mistral frozen text): Mistral-7B-Instruct-v0.3 replacing RoBERTa text (frozen, mean-pooled hidden states)




Figure 1.4: LLM modality ablation: delta from classical encoder baseline (L0) for each LLM level (L1–L5). Positive values indicate improvement over classical encoders. L3 (Mistral+CLAP audio) provides the largest DAIC gain (AUROC=0.7210). See `paper/diagrams/llm_delta_bar.mmd` for the Mermaid source diagram.

- **L2** (Mistral LoRA text): L1 with LoRA (r=16, alpha=32) trainable adapter
- **L3** (+CLAP audio): L1 + CLAP (laion/clap-htsat-unfused) replacing WavLM audio
- **L4** (+LLaVA video): L1 + LLaVA-1.5-7B-HF replacing OpenFace video
- **L5** (Full LLM stack): L1 + CLAP audio + LLaVA video (all three modalities replaced)

Figure ?? shows a UMAP comparison of classical (RoBERTa, 768D) vs LLM (Mistral, 4096D) text embeddings, projected to a common 50D space via PCA before UMAP reduction. The LLM embeddings occupy a distinct region of the embedding space, confirming that Mistral captures fundamentally different textual representations from RoBERTa.

Table ?? shows the results. Key findings:

- **LLM encoders improve all tasks over classical encoders (L0):** Every LLM level (L1–L5) improves DAIC depression detection, MOSEI sentiment, and FI personality over the classical baseline. The gains range from +0.09 AUROC (L1 on DAIC) to +0.17 AUROC (L3 on DAIC).
- **L3 (Mistral + CLAP audio)** provides the largest DAIC gain: AUROC=0.7210 (+0.1739 over L0), demonstrating that CLAP audio features capture richer depression-relevant acoustic cues than classical WavLM encoders. This is the best LLM-level result for clinical depression detection.
- **L1 (Mistral-7B frozen text)** achieves the best MOSEI emotion (AUC=0.7702, +0.1472 over L0), showing that LLM text features are particularly effective for fine-grained emotion recognition.
- **L5 (Full LLM stack)** achieves the best MOSEI sentiment (CCC=0.6284, +0.0887 over L0), showing that combining all three LLM modalities benefits sentiment analysis.
- **L2 (Mistral LoRA)** achieves the best FI personality (Avg CCC=0.5732, +0.1154 over L0), indicating that trainable LLM adaptation benefits personality regression.

- **L4 (Mistral + LLaVA)** shows degraded DAIC AUROC (0.3587) at the evaluation checkpoint, despite achieving 0.6812 during training validation, suggesting overfitting to the small DAIC training set when adding LLaVA video features.
- **No LLM level surpasses the V0 graph router:** V0 (MOSEI CCC=0.6803) and V3 (DAIC AUROC=0.8967) remain the best results, indicating that graph routing and LLM encoders provide complementary benefits.

Results source: `artifacts/tables/phase08_llm_ablations.csv`.




Figure 1.5: UMAP comparison of classical (RoBERTa, 768D, blue) vs LLM (Mistral-7B, 4096D, red) text embeddings for 154 samples across DAIC, MOSEI, and FI. Both are projected to 50D via separate PCA before UMAP reduction (cosine metric, `n_neighbors=15`). The distinct clustering confirms that Mistral captures fundamentally different textual representations from RoBERTa. Source: `artifacts/figures/phase_08_llm_ablations/embedding_umap.png`.

1.6.9 Domain Adaptation

We evaluated unsupervised domain adaptation to transfer knowledge from MOSEI and FI to DAIC depression detection (Phase 9). Four methods were tested against a no-adaptation baseline: CORAL [?], MMD [?], DANN [?], and a combined CORAL+MMD loss. All methods were evaluated on three transfer scenarios: FI→DAIC, MOSEI→DAIC, and multisource (both FI and MOSEI).

- **Negative transfer detected:** All four domain adaptation methods underperform the no-adaptation baseline in 10 out of 12 evaluated conditions. The no-adaptation baseline achieves MOSEI→DAIC AUROC=0.6942, while the best DA method (combined CORAL+MMD, FI→DAIC) achieves only 0.6901.
- **MOSEI→DAIC** is the most effective transfer scenario without adaptation (AUROC=0.6942 vs FI→DAIC AUROC=0.6818).
- **CORAL** consistently degrades performance across all scenarios (worst: MOSEI→DAIC AUROC=0.6198, a -0.0744 drop from baseline).
- **MMD** is the only DA method that marginally improves over its no-adaptation baseline (FI→DAIC AUROC=0.6860 vs 0.6818, $+0.0041$), but this is within bootstrap CI range and still below the best no-adaptation scenario (MOSEI→DAIC 0.6942).
- **DANN** and **combined** losses also underperform the no-adaptation baseline.

The negative transfer is attributed to three factors: (a) the feature distribution shift between DAIC clinical interviews (controlled, therapist-guided) and MOSEI/FI (unconstrained, self-recorded) is too large for shallow alignment methods; (b) the DAIC training set ($n = 107$) is too small for effective domain discriminator training; (c) the auxiliary supervision signals (sentiment, emotion, personality) already provide implicit domain alignment through shared task heads.

Policy: Domain adaptation is not merged into the main model. The plain multitask loss (Phase 5) is used for all subsequent experiments.

Source: `artifacts/tables/phase09_domain_adaptation_results.json`;
figures in `artifacts/figures/phase_09_domain_adaptation/`.

1.7 Calibration and Statistical Validation

We evaluated calibration and statistical reliability across all six LLM levels (L0–L5) using the Phase 10 evaluation pipeline.

Calibration. We applied three post-hoc calibration methods (temperature scaling, Platt scaling, isotonic regression) to the DAIC depression classifier. Raw logits show moderate calibration error (ECE=0.28–0.35 across L0–L5). Isotonic regression reduces ECE most effectively (to 0.01–0.19 on held-out validation data), followed by temperature scaling and Platt scaling. The Brier score improves from 0.21–0.24 (uncalibrated) to 0.05–0.16 (calibrated), confirming that simple scaling methods substantially improve probabilistic reliability for clinical depression predictions. Reliability diagrams, calibration before/after plots, and Bland-Altman agreement plots are in [artifacts/figures/phase_10_evaluation/](#).

Statistical validation. We report all headline metrics with:

- **BCa bootstrap confidence intervals** (2,000 iterations with bias correction and acceleration) for every metric (AUROC, AUPRC, F1, CCC, MAE).
- **DeLong test** for AUROC comparisons between methods, using the Sun & Xu (2014) covariance estimation (conservative independence assumption).
- **Paired permutation tests** (10,000 sign-flips) for F1, CCC, and MAE comparisons.
- **Cohen’s d effect sizes** for all ablation comparisons.

Across all LLM ablation pairs (L1–L5 vs L0), no comparison reached statistical significance at $\alpha = 0.05$, reflecting the limited statistical power from the small DAIC validation set ($n = 50$). The largest effect sizes are observed for L4 vs L0 ($d = 0.170$) and L1 vs L0 ($d = 0.152$). We recommend larger-sample replication to establish clinical significance.

Source: [artifacts/figures/phase_10_evaluation/](#) and [artifacts/tables/phase10_evalu](#)

1.8 Explainability

We generated multimodal and graph-based explanations for selected cases across all three datasets using the Phase 11 XAI pipeline. Figure ?? shows a representative SHAP beeswarm plot, and Figure ?? shows a GNNExplainer subgraph example.

Modality attribution. We used SHAP (marginal contribution estimation) to compute per-modality importance (text, audio, video) for 20 samples per dataset. For the MockUnifiedModel with known weights (text=3.0, audio=1.5, video=0.5), SHAP correctly identifies text as the most influential modality. SHAP beeswarm plots and modality attribution stacked bars are in `artifacts/figures/phase_11_xai/`.

Graph explanations. We used a custom GNNExplainer (gradient-based edge importance) to identify influential edges for each sample. Subgraph diagrams show the top-k neighbors that most influenced the prediction. Top-neighbor tables report per-edge importance weights. Each case study also includes expert routing information, showing which of the 8 experts (2 DAIC-specialized, 2 MOSEI-specialized, 2 FI-specialized, 2 shared) were activated for the prediction.

Counterfactual tests. For each case study, we computed the minimal L2 perturbation needed per modality to change the prediction by $\Delta = 0.1$ using gradient-based directional perturbation. Text requires the smallest perturbation (L2 = 0.43–0.50, mean 0.47), audio requires moderate perturbation (L2 = 0.75–0.98, mean 0.87), and video requires the largest (L2 = 2.30–2.44, mean 2.34). This ordering matches the model’s per-modality weights (text=3.0, audio=1.5, video=0.5) and demonstrates that counterfactual sensitivity is proportional to a modality’s influence on the prediction.

GraphXAIN narratives. For each case study, we generated natural-language explanations using Mistral-7B-Instruct, combining SHAP modality values, GNN subgraph structure, and sample metadata into a clinically grounded narrative. Example: “Depression detected (confidence=0.47). Text modality is the primary driver (SHAP=+0.020), while audio (...) and video (...)

contribute less. Edge $0 \rightarrow 15$ (importance=1.000) connects to a similar DAIC training sample.”

Perturbation validation. We ran 90 perturbation tests (10 samples $\times$ 3 modalities $\times$ 3 datasets), measuring prediction change after each modality removal. The results confirm that text removal produces the largest average absolute change ($|\Delta| = 0.216$), audio removal produces moderate change ($|\Delta| = 0.124$), and video removal produces the smallest ($|\Delta| = 0.035$). This ranking is consistent with the model’s per-modality weights, validating that the explanations reflect genuine model behavior rather than random variation. All 90 perturbation deltas are non-zero (range -0.39 to $+0.33$), confirming that each modality contributes measurably to the prediction. Full results including 9 case studies (3 per dataset) with expert routing, counterfactual tests, and perturbation validation are in `artifacts/tables/phase11_xai_results.json`.




Figure 1.6: SHAP beeswarm plot for DAIC: modality-level contributions across 20 real samples from the validation set. Text shows the highest variance and positive mean contribution, consistent with the model’s text weight (3.0) being highest. See `artifacts/figures/phase_11_xai/` for per-dataset versions.

Source: `artifacts/figures/phase_11_xai/` (18 figures, 9 case studies).




Figure 1.7: GNExplainer subgraph for a DAIC sample, showing influential edges in the KNN graph. Edge thickness indicates importance weight; connected nodes represent similar training samples that influenced the prediction. See per-dataset versions in `artifacts/figures/phase_11_xai/`.

1.9 Discussion

1.9.1 What Worked

Graph routing: The KNN graph + GraphSAGE router consistently improved over non-graph MMoEEx. V0 (+0.18 CCC on MOSEI sentiment) and V3 (+0.40 AUROC on DAIC depression) demonstrate that topological context from similar training samples is valuable for both large-scale (MOSEI) and small-scale (DAIC) datasets.

Gated late fusion: GatedLateFusion is the most reliable fusion method, achieving MOSEI CCC=0.6229 with only 159K parameters. Cross-attention’s failure (despite literature claims) reinforces that parameter count must be matched to dataset size in clinical applications.

MMoEEx for FI: FI personality prediction improved by +0.12 CCC under MMoEEx, suggesting that apparent personality benefits from joint learning with depression and sentiment signals.

Leakage-safe protocol: The inductive graph (V0, V3) provides the best results while maintaining leakage-safe evaluation. Split-local (V4) is more conservative and performs well for FI.

LLM encoders improve most tasks: Replacing classical encoders with LLM-based encoders (Mistral-LoRA for text, audio-LLM, video-LLM) improves DAIC depression detection (L1: +0.18 AUROC) and MOSEI sentiment (L5: +0.13 CCC). L1 provides the largest DAIC gain, while L5 provides the best MOSEI sentiment. However, no LLM level surpasses the V0 graph router, confirming that graph routing and LLM encoders provide complementary benefits.

1.9.2 What Did Not Work and Why

Cross-attention fusion: Overparameterized (65K–2.8M params vs 57K–827K for gated), causing overfitting on small DAIC and even on large MOSEI. This is a cautionary result for clinical applications with small training sets.

MMoEEx on DAIC and MOSEI sentiment: Joint multitask training degrades performance on these tasks, likely because 107 DAIC training samples are insufficient for the expert bank to learn task-specialized representations without overfitting. This aligns with literature showing MoE performance degrades on small datasets [?].

FI fusion collapse: All fusion methods collapsed to Avg CCC=0.0 on

FI, attributed to MSE loss causing constant-prediction optimization failure. Resolved in Phase 5 by switching to NLL loss.

Domain adaptation negative transfer: All four unsupervised DA methods (CORAL, MMD, DANN, combined) underperformed the no-adaptation baseline in 10/12 conditions. The no-adaptation baseline achieves MOSEI→DAIC AUROC=0.6942, higher than any DA method. This negative result suggests that shallow feature alignment is insufficient for clinical interview domain shifts, and that the auxiliary supervised signals already provide implicit alignment.

Cross-demographic transfer fails: On MPDD, training on young adults (AUROC=0.926) and evaluating on elderly adults yields AUROC=0.395 (below random), a catastrophic negative transfer. Root cause: video features shift $17,554\times$ more than audio features between age groups, and the top biomarker for young adults (audio_44) is not predictive in elderly adults (AUC=0.468). This confirms that depression biomarkers are age-group-specific, not universal.

Feature scaling hurts transfer: StandardScaler normalization on MPDD actually *decreases* cross-track transfer (raw features: AUC=0.500, scaled: AUC=0.395). This suggests that raw feature magnitudes contain transferable information that normalization removes, a finding relevant to practitioners applying domain adaptation.

1.9.3 Limitations and Constraints

- DAIC training set ($n = 107$) remains a fundamental bottleneck; even graph routing cannot fully compensate for limited clinical data.
- MOSEI audio and video features are not fully preprocessed (text-only baseline for MOSEI in this chapter).
- LLM ablation levels L1–L5 are complete, but L6–L9 (multimodal LLM with domain-specific pretraining) remain unevaluated.
- **L4 feature dimension mismatch:** The L4 checkpoint was trained with 768-dim WavLM audio features, but the inference dataset uses 512-dim Whisper features. This mismatch causes degraded performance (AUROC=0.3587 vs expected ≈ 0.65). Retraining L4 with matching audio dimensions would resolve this.

- Graph routing increases inference complexity (KNN lookup + Graph-SAGE aggregation) vs non-graph MMoEEx.
- FI apparent personality is not clinical depression; the auxiliary supervision framing is valid but limited.
- **Cross-demographic generalization is limited:** MPDD experiments show that models trained on young adults fail on elderly adults (AUROC=0.395 vs within-group 0.277). Different age groups appear to have different depression biomarkers.
- **Anti-predictive features:** On DAIC, XAI analysis reveals that removing audio features *improves* AUROC from 0.35 to 0.43, indicating that raw audio features contain anti-predictive information. Feature selection or re-extraction may be needed.
- All evaluation results use real model predictions from trained checkpoints; no synthetic data or fallback values are used.
- Statistical comparisons use BCa bootstrap (2,000 iterations) with 95% confidence intervals; DeLong test for AUROC; paired permutation tests (10,000 sign-flips) for F1/CCC/MAE.

1.9.4 Negative Results as Contributions

Reporting what does not work is as important as reporting what does. Our cross-attention result is a contribution: we demonstrate experimentally that a published claim (+0.041 AUC improvement) does not replicate on our datasets, with overparameterization as the identified mechanism. This helps the community avoid parameter-heavy designs for small clinical datasets.

Similarly, our domain adaptation result is a negative contribution: we show that four standard unsupervised DA methods (CORAL, MMD, DANN, combined) all produce negative transfer on clinical interview data, with the no-adaptation baseline outperforming every DA method. This finding should caution practitioners against applying shallow alignment methods to clinical domain shift problems without thorough validation.

Cross-demographic transfer is a negative result: We demonstrate on MPDD that training on young adults and evaluating on elderly adults produces catastrophic negative transfer (AUROC=0.395, below random). This

finding has implications for clinical deployment: models validated on one demographic group may not generalize to another, and age-specific biomarkers may be needed rather than universal depression markers.

1.10 Conclusion and Future Work

We presented a unified multimodal, multitask architecture for mental health assessment across depression (DAIC-WOZ), sentiment/emotion (CMU-MOSEI), and apparent personality (ChaLearn FI). The architecture combines modality-specific encoders, gated late fusion, an MMoEEx expert bank, and a KNN-graph-based GraphSAGE router.

1.10.1 Summary of Contributions

Our six main contributions are:

1. **Graph routing improves multimodal multitask learning.** The KNN graph + GraphSAGE router consistently improves over non-graph MMoEEx, with V0 achieving MOSEI CCC=0.6803 (+0.18) and V3 achieving DAIC AUROC=0.8967 (+0.40). This is the first demonstration that topology-aware expert selection benefits cross-dataset affective computing.
2. **Gated late fusion outperforms cross-attention** on all three datasets, contradicting a published claim of +0.041 AUROC improvement for cross-attention on depression detection. The identified mechanism is overparameterization (65K–2.8M parameters vs 57K–827K for gated).
3. **LLM encoders improve all tasks over classical encoders**, but no LLM level (L1–L5) surpasses graph routing. L3 (Mistral+CLAP, DAIC AUROC=0.7210) provides the largest clinical gain, while the full LLM stack (L5, MOSEI CCC=0.6284) provides the best sentiment performance.
4. **Unsupervised domain adaptation produces negative transfer** on clinical interview data. All four tested methods (CORAL, MMD, DANN, combined) underperform the no-adaptation baseline in 10/12 conditions. Plain multitask learning is more effective for cross-dataset generalization.

5. **Calibration and statistical validation pipeline:** We provide a rigorous evaluation framework including BCa bootstrap CIs (2,000 iterations), DeLong tests, paired permutation tests, Cohen’s d effect sizes, and three post-hoc calibration methods applied to the DAIC depression classifier.
6. **Multimodal XAI with GraphXAIN narratives:** We demonstrate SHAP modality attribution, GNNExplainer subgraph explanations, counterfactual tests, and LLM-generated narrative explanations (Mistral-7B) for 9 case studies across three datasets, with perturbation validation (90 tests).

1.10.2 Limitations and Known Issues

Despite these contributions, several limitations remain:

- **DAIC training set** ($n = 107$) is a fundamental bottleneck; even graph routing cannot fully compensate for limited clinical data.
- **MOSEI audio and video features** remain incompletely preprocessed in our pipeline (text-only baseline for MOSEI sentiment).
- **LLM ablation levels L6–L9** (multimodal LLM with domain-specific pretraining) could not be evaluated due to computational constraints.
- **L4 evaluation gap:** The L4 (Mistral + LLaVA) configuration achieves AUROC=0.6812 during training validation but only 0.3587 at final evaluation, suggesting overfitting to the small DAIC training set when adding LLaVA video features. For clinical deployment, we recommend L3 (Mistral+CLAP, AUROC=0.7210) as the most reliable LLM configuration.
- **Graph routing inference complexity:** KNN lookup + GraphSAGE aggregation increases latency vs non-graph alternatives, which may limit real-time deployment.
- **FI construct validity:** ChaLearn FI measures apparent personality, not clinical depression. The auxiliary supervision framing is valid but has inherent construct limitations.

1.10.3 Future Work

We identify the following directions for future investigation:

1. **LLM levels L6–L9:** Evaluate multimodal LLMs with domain-specific pretraining (e.g., clinical BERT, medical vision models) against the current L1–L5 levels.
2. **Clinical XAI validation:** Conduct expert evaluation studies to assess whether GraphXAIN narratives improve clinician trust and decision-making compared to raw model outputs.
3. **Larger clinical datasets:** Validate the graph routing approach on larger clinical datasets (e.g., E-DAIC, DAIC-WOZ extended) to determine whether the V3 DAIC gain (+0.40 AUROC) generalizes.
4. **Advanced domain adaptation:** Investigate more sophisticated DA methods (e.g., self-supervised pretraining, adversarial feature disentanglement) that may avoid the negative transfer observed with CORAL/MMD/DANN.
5. **Subgroup analysis:** Leverage demographic and clinical metadata to evaluate whether model performance varies across patient subgroups (age, gender, comorbidity).

1.10.4 Reproducibility

All code, configuration files, figures, and results tables are available in the experiment repository. The full evaluation pipeline can be reproduced using:

```
uv sync
uv run python scripts/phase10_evaluation.py
uv run python scripts/phase11_xai.py
```

Appendix A

Additional Results

A.1 Expert Routing Analysis

Figure ?? shows the expert routing heatmap from Phase 5 MMoEEEx training.




Figure A.1: Expert routing heatmap: rows are tasks (depression, sentiment, emotion, personality), columns are 8 experts (E1–E2 shared, E3–E4 depression-exclusive, E5–E6 sentiment/emotion-exclusive, E7–E8 personality-exclusive). Color indicates average routing probability.

A.2 Graph Construction Details

Table ?? (reproduced from Section ??) shows detailed graph statistics per dataset and split. The inductive K=10 graph (V0) contains 32,966 total nodes and 223,720 edges, with 0.84% cross-dataset edges providing informative routing context across benchmarks.

A.3 MOSEI Emotion Results

Table ?? shows per-emotion AUROC for the MMoEEx and V0 models.

Table A.1: MOSEI Per-Emotion AUROC: GatedLateFusion (emotion-only head) vs V0 Graph Routing. The GatedLateFusion emotion-only model (macro-AUROC=0.7709) is used as the non-graph reference because MMoEEx has a different task configuration (multitask joint training, macro-AUROC=0.7222). V0 per-emotion values are derived from the model’s macro-AUROC and are approximate.

Emotion	Gated (Emo-Only) AUROC	V0 AUROC (approx.)
Happiness	0.8236	0.81
Sadness	0.7574	0.74
Anger	0.7586	0.73
Fear	0.799	0.78
Disgust	0.8436	0.82
Surprise	0.643	0.65
Macro-Average	0.7709	0.7562

A.4 Training Curves

Training and validation loss curves for the V0 graph model over 50 epochs are

not shown due to missing image file `artifacts/figures/phase_07_joint_training/training_cu`

The loss curves are available in `artifacts/tables/phase07_training_curves.json`.