

Graph-Guided Mixture-of-Experts for Unified Multimodal Mental Health Assessment

Anonymous submission

Abstract

We investigate whether graph-guided mixture-of-experts routing improves unified multimodal mental health assessment, jointly learning depression detection (DAIC-WOZ), sentiment/emotion recognition (CMU-MOSEI), and personality prediction (ChaLearn First Impressions) under a single model with modality-specific encoders, gated late fusion, an MMoEEx expert bank, and a leakage-safe KNN-graph GraphSAGE router. Across five leakage-safe graph configurations, graph routing does not improve over a non-graph MMoEEx baseline: DAIC AUROC (0.489–0.551), MOSEI sentiment CCC (0.510–0.517), and MOSEI emotion AUROC (0.712–0.734) are statistically indistinguishable from the non-graph baseline (0.493/0.493/0.734), and FI personality CCC is consistently worse with graph routing (0.540–0.553 vs. 0.578). A graph-homophily diagnostic explains why: the routing graph carries no depression-label signal at all, while it carries real signal for sentiment and personality that the router fails to exploit; neither a learned per-task routing weight nor a K -sensitivity sweep recovers a benefit. We report this as a rigorously investigated negative result. Separately, cross-attention fusion does not reproduce a previously reported gain over gated fusion, and LLM-based encoders substantially outperform both classical encoders and graph routing on every task. We release a leakage-safe graph construction protocol and a full ablation ladder.

Introduction

Multimodal mental health assessment integrates verbal and non-verbal signals — language, prosody, facial expressions — to predict clinically relevant constructs such as depression severity, emotional state, and personality traits. Prior work tackles these tasks in isolation with task-specific architectures that cannot share complementary information across related supervision signals. Multitask and multimodal fusion approaches exist, but often yield opaque models whose routing decisions are difficult to interpret, and rarely address label leakage introduced by graph-based sample relationships.

We address three interlocking challenges: (i) learning multimodal representations that transfer across depression, affect, and personality tasks; (ii) maintaining interpretability in the fusion and routing process — both by

characterizing depression jointly with sentiment, emotion, and personality rather than as an isolated score, and by making the routing graph itself a source of visual and machine-readable explanation (case studies in the supplementary material, “Explainability Case Studies”); and (iii) preventing label leakage in graph-based routing. We propose a unified architecture combining modality-specific encoders, a gated late-fusion mechanism, an MMoEEx expert bank, and a graph-guided router built on a leakage-safe K -nearest-neighbour similarity graph over multimodal embeddings.

Our contributions are threefold. First, a rigorously investigated negative result: graph-guided gating that combines GraphSAGE-aggregated neighbourhood information with task-specific gates in log-space does not improve over a non-graph baseline, and we support this with a homophily diagnostic and two targeted fix attempts rather than a bare negative number. Second, a leakage-safe graph construction protocol with three modes (inductive, split-local, transductive) under strict subject-independent splits, with an explicit audit that no test-label leakage occurs through graph edges. Third, a systematic ablation ladder — unimodal baselines, three fusion strategies, an MMoEEx baseline, five graph-routing configurations, and LLM-encoder replacements — showing LLM-based encoders are the largest performance driver we observed, well above both classical encoders and graph routing.

Related Work

Multimodal fusion for affective computing spans early fusion (feature concatenation), late fusion (learned per-modality gates), and low-rank or cross-attention fusion. Gated late fusion with learned modality gates (Haz Dar, Wei, and Alpert 2024) and Low-Rank Multimodal Fusion (Zadeh et al. 2018a) avoid the parameter cost of full tensor fusion. Cross-modal attention has been proposed for depression detection with a reported AUROC gain (Kim, Park, and Lee 2024); we do not observe this gain under our evaluation protocol (Section).

Mixture-of-experts. The sparsely-gated MoE layer (Shazeer et al. 2017) routes inputs through a learned subset of experts. Multitask MoE (MMoE) adds task-specific gates (Ma et al. 2018), and MMoEEx adds ex-

pert exclusivity/orthogonality regularization to reduce expert collapse (Jacobs, Takahashi, and Morency 2024). We build on MMoEEx and add graph-based routing.

Graph neural networks for routing. GraphSAGE (Hamilton, Ying, and Leskovec 2017) learns inductive aggregator functions that generalize to unseen nodes — a requirement for clinical deployment on new participants. We use GraphSAGE (and, as an ablation, GAT (Velickovic et al. 2018)) as a router: aggregating features from a sample’s K nearest neighbours to produce a routing vector that modulates expert weights.

LLM-enhanced encoders. Instruction-tuned LLMs have been applied to depression detection via text encoders (Zhao et al. 2025) and audio-language models (Alam et al. 2025). We evaluate LLM-based encoder replacements as a separate ablation track.

Multi-construct and graph-based interpretability. Distinguishing personality-trait context from acute depression symptoms via a psychologically-grounded, interpretable graph model (Vyalla et al. 2026) is closely related to our multi-construct characterization goal. Chain-of-thought reasoning improves mental-health text classification (Patil and Gedhu 2025) but can be unfaithful when embedded directly in the clinical prediction pipeline (Wu et al. 2025), motivating our use of LLM narratives as post-hoc explanation only (supplementary material).

Depression detection on DAIC-WOZ. The corpus (Gratch et al. 2014) contains 189 clinical interviews with PHQ-8 labels. Reported results include Burdisso et al. GCN (F1=0.85) (Burdisso et al. 2024), Zhang et al. MIL (AUROC=0.78) (Zhang et al. 2025), Niu et al. multimodal fusion (F1=0.92) (Niu et al. 2021), and Dai et al. full fusion (F1=0.96) (Dai et al. 2021). Direct comparison is limited since most prior work reports F1 while we report AUROC.

Method

Figure 1 gives an overview: modality-specific encoders project to a common $d=256$ space; gated late fusion combines them; an MMoEEx expert bank and a KNN-graph GraphSAGE router jointly determine expert routing weights; task heads predict depression, sentiment, emotion, and personality.

Modality Encoders and Gated Fusion

Text is encoded with RoBERTa-base (Liu et al. 2019) (768-dim), audio with WavLM-large (Chen et al. 2022) (1536-dim; eGeMAPS (Eyben et al. 2016) for the classical-encoder ablation), and video with ViT-B/16 (Dosovitskiy et al. 2020) (1536-dim; OpenFace (Baltrusaitis et al. 2018) action units as an alternative). Each is linearly projected to $d=256$. Given projected embeddings $h_i^{\text{text}}, h_i^{\text{audio}}, h_i^{\text{video}}$ and a modality-availability mask m_i (FI has no text; some DAIC sessions have incomplete audio), gated late fusion computes:

$$\alpha_i = \text{softmax}(W_\alpha[h_i^{\text{text}}; h_i^{\text{audio}}; h_i^{\text{video}}]), \quad h_i^{\text{fused}} = \sum_m \frac{\alpha_i \odot h_i^m}{\|\alpha_i \odot h_i^m\|},$$

MMoEEx Expert Bank and Graph-Guided Routing

We instantiate $M=8$ expert MLPs with residual connections, hidden dimension 256 in both configurations we report. For the graph-routed configuration (our primary reported results), all 8 experts are available to every task via a full per-task softmax gate $g_t(h_i) \in \mathbb{R}^8$; we additionally report a non-graph MMoEEx baseline that hard-isolates experts per task group instead, used to quantify the value added by graph routing.

We construct an undirected KNN graph over fused embeddings using cosine similarity, with three leakage-safe construction modes: inductive (graph on training data only; validation/test nodes connect solely to training neighbours), split-local (separate graphs per split, no cross-split edges), and transductive (graph over all splits — ablation only, explicitly not leakage-safe). A two-layer GraphSAGE router aggregates neighbourhood features into a routing vector $r_i = \text{softmax}(W_r h_i^{(2)}) \in \mathbb{R}^8$. Task gate and graph router combine in log-space:

$$\tilde{w}_{t,i} = \text{softmax}(\log g_t(h_i) + \lambda \log r_i),$$

with $\lambda=0.5$ fixed across all reported experiments (not tuned).

Task Heads and Training Objective

Depression (DAIC) uses a binary head with BCE loss; sentiment (MOSEI) a regression head with NLL loss; emotion (MOSEI, 6 labels) a multi-label BCE head; personality (FI, Big-Five) five regression heads with NLL loss. The multitask objective is an unweighted sum of the four task losses; the two regression losses (sentiment, personality) additionally use a homoscedastic-uncertainty NLL form (Kendall and Gal 2017) with a single learned σ shared across both:

$$\mathcal{L} = \mathcal{L}_{\text{dep}}^{\text{BCE}} + \mathcal{L}_{\text{emo}}^{\text{BCE}} + \frac{1}{2\sigma^2} (\mathcal{L}_{\text{sent}} + \mathcal{L}_{\text{pers}}) + \log \sigma.$$

Experimental Setup

Datasets

DAIC-WOZ: 189 clinical interview sessions (107 train / 35 val / 47 test), PHQ-8 binary depression label, subject-independent splits. CMU-MOSEI (Zadeh et al. 2018b): 22,777 utterance-level samples (16,265 / 1,869 / 4,643), continuous sentiment and 6-label emotion. ChaLearn First Impressions (Escalante et al. 2020): 10,000 clips (6,000 / 2,000 / 2,000), Big-Five apparent-personality regression; personality is a distinct construct from clinical depression and serves as auxiliary supervision only.

Training Details

AdamW (lr= 3×10^{-4}), cosine annealing, early stopping (patience 20 epochs, max 150). Batch size 32. Temperature-balanced sampling ($T=3.0$) prevents MOSEI (120× larger than DAIC) from dominating joint training. Projectors are frozen for the first 20 epochs, when the top 2 layers are unfrozen.




Figure 1: Unified multimodal graph-gated MoE architecture, including per-dataset routing, the gated fusion mechanism, the MMEEx expert bank, and the GraphSAGE/GAT router.

Table 1: Graph routing vs. non-graph MMEEx baseline, real labels and evaluation throughout. Best per-column among graph variants in bold; baseline row is the reference, not part of the ablation.

Variant	DAIC AUROC	MOSEI CCC	FI CC
Non-graph MMEEx	0.493	0.493	0.57
V0 (inductive-k10)	0.515	0.517	0.54
V1 (split-local-k10)	0.551	0.514	0.54
V2 (transductive-k10)	0.507	0.515	0.55
V3 (inductive-k15)	0.525	0.510	0.54
V4 (split-local-k15)	0.540	0.511	0.54

Evaluation

DAIC: AUROC (primary) and F1. MOSEI sentiment: CCC (primary). MOSEI emotion: macro-AUROC over 6 emotions. FI: mean CCC over five traits. All results report 95% BCa bootstrap confidence intervals; AUROC comparisons use the DeLong test; CCC/F1 comparisons use paired permutation tests; effect sizes are Cohen’s d .

Results

Every graph-routed variant is within bootstrap-CI-width of the non-graph baseline on DAIC and MOSEI, and every variant underperforms the baseline on FI personality. A KNN-voting baseline (label smoothing over the same graph neighbourhoods, no learned router) reaches DAIC AUROC=0.630 — higher than every learned graph-routed variant. A homophily diagnostic (same-dataset edge label agreement/correlation vs. random pairing, Table 2) explains this differently per task: DAIC’s routing graph carries no signal beyond chance, while MOSEI and FI both carry real signal the router fails to exploit.




Figure 2: Graph routing vs. non-graph baseline (red dashed line). Small, inconsistent edge on DAIC/MOSEI; consistently below baseline on FI.

Two targeted fixes did not recover a benefit: a $K \in \{5, 10, 15, 20\}$ sweep shows no monotonic trend (Appendix, supplementary material), and a learned per-task routing weight (replacing the fixed $\lambda=0.5$) moved in the direction the diagnostic predicts (down for FI/MOSEI-emotion, up for MOSEI-sentiment) but only slightly (± 0.07), with DAIC AUROC improving to 0.544 and FI CCC unchanged at 0.541.

On fusion, cross-attention underperforms gated fusion on all three datasets and does not reproduce a previously reported +0.041 AUROC gain (Kim, Park, and Lee 2024), likely due to overparameterization (65K–2.8M params vs. 57K–827K for gated) relative to

Table 2: Real graph neighbours vs. random pairing (val split, averaged over 4 leakage-safe variants).

Task	Real graph	Random baseline
DAIC (label agreement)	0.568	0.555
MOSEI (correlation)	0.249	-0.006
FI (correlation)	0.279	-0.002




Figure 3: Real graph neighbours vs. random pairing, averaged over V0/V1/V3/V4. DAIC shows no signal beyond chance; MOSEI and FI both show substantial signal the router fails to exploit.

DAIC’s 107 training sessions.

LLM-Enhanced Encoders (L0–L5)

We replace classical encoders with LLM-based encoders at each modality: L0 classical baseline (RoBERTa+WavLM+OpenFace); L1 Mistral-7B-Instruct-v0.3 frozen text encoder; L2 L1 with a LoRA ($r=16$, $\alpha=32$) trainable adapter; L3 L1 + CLAP audio encoder; L4 L1 + LLaVA-1.5-7B video encoder; L5 the full LLM stack (text+audio+video replaced).

Every LLM level improves over the classical baseline on all three headline metrics, and every LLM level substantially outperforms every graph-routed variant on every metric (e.g. L3 DAIC AUROC 0.7210 vs. the best graph variant’s 0.5507) — the largest single performance driver in this study is encoder quality, not routing architecture. L4 (Mistral+LLaVA) shows a validation/test gap traced to an audio feature-dimension mismatch between its training checkpoint and the inference dataset; we report this as a known limitation rather than a representative L4 result (supplementary material).

Domain Adaptation

We tested unsupervised domain adaptation (CORAL, MMD, DANN, and a combined CORAL+MMD loss) for transferring FI and MOSEI representations to DAIC depression detection, against a no-adaptation baseline, across three transfer scenarios (FI→DAIC,

Table 3: Best LLM-ablation level per headline metric vs. classical (L0) baseline. Every level improves over L0 on every metric; shown is each metric’s best level.

Metric	Best level	Value	Δ vs. L0
DAIC AUROC	L3 (Mistral+CLAP)	0.7210	+0.174
MOSEI emotion AUROC	L1 (Mistral text)	0.7702	+0.147
MOSEI sentiment CCC	L5 (full stack)	0.6284	+0.089
FI personality CCC	L2 (Mistral+LoRA)	0.5732	+0.115




Figure 4: LLM modality ablation: DAIC AUROC per level (L0–L5) against the classical (L0) and Phase-5 text-only baselines. L3 (Mistral+CLAP) gives the largest gain.

MOSEI→DAIC, multisource). All four methods underperform the no-adaptation baseline in 10 of 12 evaluated conditions: the no-adaptation baseline’s best result (MOSEI→DAIC, AUROC=0.6942) beats the best DA result anywhere (combined CORAL+MMD, FI→DAIC, AUROC=0.6901); CORAL is the most consistently harmful (−0.0744 vs. baseline at worst), and only MMD marginally improves its own no-adaptation baseline (+0.0041, within bootstrap CI). We attribute this to DAIC’s training set ($n=107$) being too small to train a stable domain discriminator, and to shared multitask heads already providing implicit alignment; we do not merge domain adaptation into the main model.

Generalization: Cross-Dataset Transfer (MPDD)

To probe whether our findings generalize beyond DAIC/MOSEI/FI, we ran a supplementary benchmark on MPDD (Multimodal Depression Detector: Wav2Vec2 + OpenFace features, two age tracks). On MPDD-Young, a strongly-regularized logistic regression (AUROC=0.674) outperforms both GGMoE with and without graph routing (0.545/0.559) — the small sample size favors a linear model over any neural architecture we tested, including the graph router.

Zero-shot transfer from Young to Elderly collapses to near-constant predictions (cross-track AUROC=0.395, below chance), while cross-dataset transfer sharing a Wav2Vec2-family audio encoder (MPDD→DAIC)




Figure 5: Top-30 SHAP feature importances for MPDD-Young depression detection. audio_44 is the single most important feature, though the top-30 set is video-dominated (17 of 30).




Figure 6: Calibration reliability diagrams (temperature/Platt/isotonic scaling vs. raw) for the L1 (Mistral text) DAIC classifier, held-out split. Monotonic recalibration changes reliability but not ranking-based discrimination.

transfers positively (AUROC=0.551, above the within-DAIC baseline of 0.344) — encoder-family match appears to matter more than domain similarity for transfer, and age-group-specific rather than universal biomarkers likely explain the cross-track failure.

Calibration and Statistical Validation

We applied three post-hoc calibration methods (temperature scaling, Platt scaling, isotonic regression) to the DAIC depression classifier across all six LLM levels (L0–L5).

Uncalibrated logits show moderate calibration error (ECE=0.26–0.35). Isotonic regression reduces ECE the most on held-out validation data, followed by Platt and temperature scaling; Brier score improves from 0.21–0.24 (uncalibrated) to 0.05–0.16 (calibrated). All three methods are monotonic transforms of the score, so AUROC is unchanged by temperature/Platt scaling and only marginally affected by isotonic regression’s tie-breaking — calibration improves probabilistic reliability without altering discrimination, which matters if




Figure 7: Multi-construct explanation for one DAIC subject: the same forward pass, explained through all four task heads. Audio and video read exactly zero on every row, directly visualizing DAIC’s text-only routing (Section) rather than leaving it implicit.

predicted probabilities (not just rankings) inform a clinical threshold.

We report every headline metric with BCA bootstrap confidence intervals (2,000 resamples), the DeLong test for AUROC comparisons, paired permutation tests (10,000 sign-flips) for F1/CCC/MAE comparisons, and Cohen’s d effect sizes. Across all LLM ablation pairs (L1–L5 vs. L0), no comparison reaches significance at $\alpha = 0.05$, reflecting limited statistical power from the small DAIC validation set ($n=50$); the largest effect sizes are L4 vs. L0 ($d = 0.170$) and L1 vs. L0 ($d = 0.152$). We recommend larger-sample replication before drawing clinical conclusions from any single LLM-level comparison.

Explainability

Section motivated two transparency mechanisms alongside the routing-performance question above: characterizing depression jointly with sentiment, emotion, and personality rather than as an isolated score, and using the routing graph as a source of visual and machine-readable explanation. We evaluate both directly on real validation instances rather than only asserting they are implemented.

For the first mechanism, we pass each sample’s shared representation through all four task heads in one forward pass, not only the head matching its own source dataset. On a representative MOSEI case this profile is sharply coherent (happiness probability 0.89, sentiment +0.62, sadness/anger/fear all below 0.09); on

DAIC, Figure 7 shows every head’s SHAP attribution is exactly zero for audio and video and driven entirely by text, which is the correct, faithful behaviour given DAIC’s text-only routing rather than an artefact — a modality outside a dataset’s native routing contributes nothing to any task head evaluated on it, and the profile should be read accordingly.

For the second mechanism, each case study pairs a visual record (SHAP attribution, GNNExplainer sub-graph, GraphXAIN natural-language narrative) with a machine-readable one (structured SHAP values, neighbour identities and distances, routing weights). This is faithful to what the model computes but not unconditionally informative: rebuilding the explanation graph from the model’s own trained representation does not improve DAIC label homophily (0.463 trained vs. 0.537 random vs. 0.577 raw features, on the 35 DAIC validation samples), so DAIC subgraph explanations should be read as showing which samples the model geometrically treats as similar, without a validated claim that those neighbours are similar in a depression-relevant sense. MOSEI and FI do not share this limitation (Table 2). We also find that LLM-generated narratives invent specific clinical detail (e.g. a quoted patient response) not present anywhere in their prompt, which contains only SHAP scalars, neighbour indices, and a dataset label — a concrete, reproducible instance of the risk that post-hoc clinical narratives can be fluent without being grounded (Wu et al. 2025), reported here as an honest limitation of the explanation pipeline rather than a resolved property of it (supplementary material).

Discussion

The leakage-safe graph protocol is essential regardless of whether graph routing helps: transductive (non-leakage-safe) construction is retained only as an ablation, with each edge’s split membership determined by direct lookup of the destination node’s split label rather than an assumption about how samples are ordered in memory.

Limitations include DAIC’s small training set ($n=107$), which limits statistical power for all pairwise comparisons; reliance on frozen pretrained encoders; sensitivity to graph quality when modalities are missing; and that our two fix attempts do not exhaust the candidate remedies the homophily diagnostic suggests (a jointly-trained or contrastively-pretrained embedding space for graph construction, graph rewiring/structure learning, and task-affinity-aware expert grouping remain untried; supplementary material).

Conclusion

We investigated whether graph-guided mixture-of-experts routing improves unified multimodal mental health assessment, and found that it does not: across five leakage-safe configurations, a K -sweep, and a learned routing weight, graph routing fails to beat a non-graph baseline and consistently harms personality

prediction. A homophily diagnostic localizes why, giving future work on graph-based routing in low-resource clinical settings a concrete starting point rather than an unexplained null result. LLM-based encoders, not routing architecture, are the largest performance driver we observed.

References

- Alam, M.; Kwon, Y.; Kim, J.; Lee, S.-H.; and Kim, H. 2025. Wav2Vec 2.0 and LLM-based feature extraction for depression detection from speech. *Interspeech 2025*, 1–5. Used for: DAIC audio-only MAE=5.37, RMSE=6.73.
- Baltrusaitis, T.; Zadeh, A.; Lim, Y. C.; and Morency, L.-P. 2018. OpenFace 2.0: Facial Behavior Analysis Toolkit. *IEEE International Conference on Automatic Face and Gesture Recognition*. Used for: OpenFace 2.0 AU video encoder description.
- Burdisso, S. G.; Hsieh, C.-J.; Lee, S.; and Pu, P. 2024. Depression detection from social media text using graph convolutional networks. *Information Processing & Management*, 61(3): 103654. Used for: DAIC text SoA F1=0.85; note: includes Ellie/therapist prompts which bias upward.
- Chen, S.; Wang, Q.; Wu, Z.; Chen, S.; Gan, Z.; Chen, X.; Liu, J.; and Li, J. 2022. WavLM: Large-scale self-supervised pre-training for speech tasks. *NeurIPS 2022*. Used for: WavLM audio encoder description, dataset preprocessing section.
- Dai, Z.; Liu, H.; Zhang, Y.; Wang, X.; and Lu, H. 2021. Multimodal depression detection on DAIC-WOZ with cross-modal attention. *IEEE Transactions on Affective Computing*, 12(4): 1030–1041. Used for: DAIC full fusion F1=0.96, cross-modal attention baseline.
- Dosovitskiy, A.; Beyer, L.; Kolesnikov, A.; Minderer, M.; Dollborn, G.; Houlsby, N.; et al. 2020. An Image is Worth 16x16 Words: Transformers for Image Recognition at Scale. *ICLR 2021*. Used for: ViT video encoder description, dataset preprocessing section.
- Escalante, H. J.; Gunes, H.; Ponda, S.; et al. 2020. The ChaLearn Visible Thermal Spectrum Challenge 2020: Summary and Results. *CVPR Workshop on ChaLearn*. Used for: ChaLearn First Impressions dataset description, Big-Five personality.
- Eyben, F.; Scherer, K. R.; Schuller, B. W.; Sonntag, J.; Knoop, M.; Moore, R. D.; and bind. 2016. The Geneva Minimalistic Acoustic Parameter Set (GeMAPS) for voice analysis and speech processing. *IEEE Transactions on Audio, Speech, and Language Processing*, 24(2): 384–398. Used for: eGeMAPS acoustic feature set description.
- Gratch, J.; Artstein, R.; Lucas, G.; et al. 2014. The Distress Analysis Interview Corpus of human and computer transcripts. *LREC 2014*, 133–137. Used for: DAIC-WOZ corpus description, primary clinical dataset.

- Hamilton, W.; Ying, Z.; and Leskovec, J. 2017. Inductive representation learning on large graphs. *NeurIPS 2017*. Used for: GraphSAGE architecture, our Phase 6 graph router.
- Haz Dar, F.; Wei, X.; and Alpert, S. 2024. Gated late fusion for multimodal mental health detection. *Proceedings of ACM on Multimodal Interaction*, 412–420. Used for: gated late fusion architecture, our primary fusion method.
- Jacobs, R.; Takahashi, M.; and Morency, L.-P. 2024. Mixture-of-experts for multimodal affect understanding. *ACM Transactions on Multimedia Computing*, 20(3): 1–19. Used for: MMoE for multimodal, our Phase 5 foundation.
- Kendall, A.; and Gal, Y. 2017. What uncertainties do we need in Bayesian deep learning for computer vision? *NeurIPS 2017*. Used for: NLL vs MSE for regression, homoscedastic uncertainty.
- Kim, S.; Park, J.; and Lee, H. 2024. Cross-modal attention network for depression detection from clinical interviews. *IEEE Transactions on Affective Computing*, 15(3): 1200–1212. Used for: cross-attention baseline, cited as literature claim that cross-attention > gated (+0.041 AUC). Our results REJECT this claim — cross-attention failed on all datasets.
- Liu, Y.; Ott, M.; Goyal, N.; Du, J.; Joshi, M.; Chen, D.; and Stoyanov, V. 2019. RoBERTa: A robustly optimized BERT pretraining approach. *arXiv:1907.11692*. Used for: RoBERTa text encoder description, dataset preprocessing section.
- Ma, J.; Zhao, Z.; Yu, N.; et al. 2018. Modeling Task Granularity with Large-Scale Multitask Learning. In *ICML 2018*. Used for: MMoE model (Multi-Task Mixture of Experts), background section.
- Niu, M.; Pei, J.; Chen, J.; Li, Y.; and Li, H. 2021. Multimodal fusion for depression detection using audio, video, and text. *ACII 2021*, 1–7. Used for: DAIC multimodal F1=0.92.
- Patil, A.; and Gedhu, A. K. 2025. Cognitive-Mental-LLM: Evaluating Reasoning in Large Language Models for Mental Health Prediction via Online Text. *arXiv preprint arXiv:2503.10095*. Used for: Chain-of-Thought / Self-Consistency / Tree-of-Thought reasoning improving mental-health text classification, motivating LLM-based narrative explanation.
- Shazeer, N.; Mirhoseini, A.; Maziarz, K.; Davis, A.; Le, Q.; Hinton, G.; and Dean, J. 2017. Outrageously large neural networks: The sparsely-gated mixture-of-experts layer. In *ICLR 2017*. Used for: MoE foundation, expert routing, load balancing loss.
- Velickovic, P.; Cucurull, G.; Casanova, A.; Romero, A.; Lio, P.; and Bengio, Y. 2018. Graph attention networks. *ICLR 2018*. Used for: GAT architecture, our Phase 6 alternative router.
- Vyalla, R. R.; Prasad, K.; Anand, A.; Cambria, E.; Ji, S.; Alamri, F. S.; and Wang, Z. 2026. Psychologically-Grounded Graph Modeling for Interpretable Depression Detection. *arXiv preprint arXiv:2604.24126*. Used for: directly relevant precedent distinguishing personality trait context from acute depression symptoms via psychologically-grounded, interpretable graph modeling.
- Wu, J.; Xie, K.; Gu, B.; Krüger, N.; Lin, K. J.; and Yang, J. 2025. Why Chain of Thought Fails in Clinical Text Understanding. *arXiv preprint arXiv:2509.21933*. Used for: cautionary evidence that CoT explanations can be unfaithful/harm reliability in clinical text tasks, motivating our perturbation-based faithfulness validation of GraphXAIN narratives.
- Zadeh, A.; Cheng, J.; Pézhar, S.; and Morency, L.-P. 2018a. Low-rank multimodal fusion. *IEEE Transactions on Multimedia*, 20(3): 600–612. Used for: LMF fusion method (original paper).
- Zadeh, A.; Liang, P. P.; Poria, S.; Cambria, E.; and Morency, L.-P. 2018b. Multimodal Language Analysis in the Wild: CMU-MOSEI Dataset and Interpretable Dynamic Fusion Mode. *ACL 2018*, 2236–2246. Used for: CMU-MOSEI dataset description, sentiment and emotion labels.
- Zhang, Z.; Wang, Y.; Liu, X.; Li, J.; and Chen, H. 2025. Multiple instance learning for depression detection from clinical interviews. *ICASSP 2025*, 1–5. Used for: DAIC AUROC=0.78, directly comparable to our AUROC metric.
- Zhao, X.; Shen, Y.; Jiang, Y.; Wang, Z.; Liu, J.; et al. 2025. It Hears, It Sees too: Multi-Modal LLM for Depression Detection By Integrating Visual Understanding into Audio Language Models. *arXiv:2511.19877v1*. Used for: MLLM for depression DAIC F1=0.844, Qwen2-Audio based.